



MPI – Message Passing Interface

CHAPTER 4


Outline

1. MPI Overview
2. Basic Structure of a MPI program
3. Messages and Point-to-Point Communication
- 4. Non-blocking Communication**
5. Derived Data Types
6. Collective Communication



Non-blocking Communication

Why nonblocking communication?



Blocking case

Non-blocking case

Blocking Communication requires processes to stop working to send and receive messages

Non Blocking Communication permits processes to keep working while checking once in a while if a message has been sent




Non-blocking Communication

Deadlock

- Code in each MPI process:

```
MPI_Ssend(...,right_rank,...)
```

```
MPI_Recv (... ,left_rank,...)
```



will block and never return because MPI_Recv cannot be called in the right-hand MPI process

- Same problem with standard send mode (MPI_Send), if MPI implementation chooses synchronous protocol



Non-blocking Communication

Separate communication into three phases:


- **Initiate nonblocking communication that returns Immediately**
(routine name starting with **MPI_I...**)

In the example: Initiate nonblocking send to the right neighbor (□)

- **Do some work**

In the example: Receiving the message from left neighbor (□)

- **Wait for nonblocking communication to complete (/)**



Non-blocking Communication

Nonblocking Send

```
MPI_Isend(void *buf, int count, MPI_Datatype datatype, int  
dest, int tag, MPI_Comm comm, MPI_Request *request)
```

```
MPI_Wait(MPI_Request *request, MPI_Status *status)
```

- *buf* must not be modified between *Isend* and *Wait*
- *Isend* + *Wait* directly after *Isend* is equivalent to blocking call (*Ssend*)
- *status* is not used.



Non-blocking Communication

Nonblocking Receive

```
MPI_Irecv (void *buf, int count, MPI_Datatype datatype, int  
source, int tag, MPI_Comm comm, MPI_Request *request)
```

```
MPI_Wait(MPI_Request *request, MPI_Status *status)
```

- **buf** must not be modified between *Irecv* and *Wait*
- Message **status** is returned in *Wait*
- **Request handle:**
 - Must be stored in local variables: **MPI_Request**
 - Is generated by a nonblocking communication routine
 - Is used (and freed) in the **MPI_WAIT** routine



Non-blocking Communication

Blocking and Non-Blocking

- Send and receive can be blocking or nonblocking
- A blocking send can be used with a nonblocking receive, and vice-versa
- Nonblocking sends can use any mode
 - standard: `MPI_ISEND`
 - synchronous: `MPI_ISSEND`
 - buffered: `MPI_IBSEND`
 - ready: `MPI_IRSEND`
- Synchronous mode affects completion, i.e. `MPI_Wait` / `MPI_Test`,



Non-blocking Communication

Completion

```
MPI_Wait(MPI_Request *request, MPI_Status *status)
```

```
MPI_Test(MPI_Request *request, int *flag, MPI_Status *status)
```

- Completion can be checked by:
 - **Wait:** It blocks the task (does not go back of the routine) until the communication has finalized.
 - **Test:** It gives back a value TRUE or FALSE depending if the communication has finalized or not.



Non-blocking Communication

Example

```
/* Blocking*/
MPI_RECV(x,N,MPI_Datatype,...,&status)

/* Non Blocking with wait */
MPI_Irecv(x,N,MPI_Datatype,...,&request)
... Realise unrelated work with the vector X
MPI_WAIT(&request, &status)
... Realise work related with the vector X


/* Non Blocking with test */
MPI_Irecv(x,N,MPI_Datatype,...,&request)
MPI_TEST(&request, &flag, &status)
while (flag==FALSE) {
    ... Realise unrelated work with the vector X
    MPI_TEST(&request, &flag, &status)
}
... Realise work related with the vector X
```



Non-blocking Communication


Activity 4: Rotating information around a ring

1. A set of processes are arranged in a ring.
2. Each process stores its rank into an integer variable *snd_buf*.
3. Each process passes this on to its neighbour on the right.
4. Each processor calculates the sum of the values.
5. Repeat steps 2-5 with “size” iterations (size = number of processes)
6. Implement the program using blocking operations and verify the correctness.
7. Use nonblocking MPI_Issend and verify the correctness.



Non-blocking Communication

Activity4: Rotating information around a ring



Thanks for your
attention!

MPI – Introduction to the Message Passing Interface

CHAPTER 4

