

Communications Services and Security **Quality of Service**

Cèsar Fernández

Departament d'Informàtica
Universitat de Lleida

Curs 2022 - 2023



1 QoS Overview

- What is QoS ?
- QoS architecture
- QoS service models
- CISCO IOS images

2 Classification

- Classification overview
- Policy-based routing
- Policy propagation via BGP
- Committed Access Rate (CAR)

3 Congestion Management

- Congestion management overview
- Flow-based WFQ
- Class-based WFQ (CBWFQ)
- Custom queueing (CQ)
- Priority queueing (PQ)

- Low Latency Queueing (LLQ)

4 Congestion Avoidance

- Congestion avoidance overview
- WRED

5 Policing and shaping

- Policing and shaping overview
- Token bucket
- Traffic policing
- Traffic shaping

6 Resource Reservation Protocol

- Overview
- The protocol
- Features
- Configuration and monitoring

7 Bibliography



Contents

- 1 **QoS Overview**
 - What is QoS ?
 - QoS architecture
 - QoS service models
 - CISCO IOS images
- 2 Classification
- 3 Congestion Management
- 4 Congestion Avoidance
- 5 Policing and shaping
- 6 Resource Reservation Protocol
- 7 Bibliography



What is QoS ?

Ability of a network to improve service to specific network traffic, providing the following services:

- Dedicated bandwidth
- Improving packet losses
- Avoiding and managing congestion
- Shaping traffic
- Setting priorities across the network



QoS architecture

Three essential components:

- QoS in a single network: queuing, scheduling and shaping
- QoS across networks: signaling
- QoS policy and management

Types of routers:


- Edge routers: packet classification, admission control
- Backbone routers: congestion management and avoidance



QoS architecture

Queuing

- Soft queues only formed when incoming traffic is faster than outgoing rates
- By default (if not QoS defined), slow output i/fs (few Mbps) use *Weighted Fair Queuing*. Otherwise: FIFO applied
- Queue length may be configured
- When queues are full, traffic is dropped



Source: [Queuing Principles](#)

QoS architecture

Scheduling


- How the soft queues are served:
 - WFQ (Flow based, class based)
 - Custom queuing: assigns a given bandwidth
 - Priority queuing: Assigns priority. Higher priorities are served first



QoS architecture

Shaping


- Average rate and maximum burst size are enforced on outgoing traffic
- Token bucket mechanisms



QoS architecture

Signaling

- Field TOS (Type of Service) of IPv4 header marked to indicate priority
- 3 MSB determine **IP precedence**. 8 priority levels
- 6 MSB determine DSCP (Differentiated Services Code Point, **DiffServ**). New standard.



IP Precedence values			
111	Network Control	011	Flash
110	Internet Control	010	Immediate
101	Critical	001	Priority
100	Flash override	000	Routine



QoS service models

Service models (or levels of service) describe the end-to-end QoS capabilities. 3 models:

- Best effort
- Integrated services
- Differentiated services

Best effort

Network delivers data if it can, without any assurance of reliability, delay bounds, or throughput.

FIFO queuing. Suitable for most applications (email, file transfer, ...)



QoS service models

Integrated services

- Application requests a specific service before sending data
- Requests made by signaling (e.g. RSVP (Reservation Protocol), asking for bandwidth and delay requirements)
- If possible, networks employs smart queuing mechanisms to provide service; WFQ or WRED (Weighted RED)

Differentiated services

- Not explicitly requested service
- Using IP Precedence or DCSP signaling

CISCO IOS images

- Images that supports QoS commands in this course:
 - IOS 12. `c7200-adventerprisek9-mz.124-24.T5`
 - IOS 15. `c7200-advipservicesk9-mz.150-1.M`
- Image `c7200-adventerprisek9-mz.152-4.M7` doesn't support some QoS commands



Contents

- 1 QoS Overview
- 2 **Classification**
 - Classification overview
 - Policy-based routing
 - Policy propagation via BGP
 - Committed Access Rate (CAR)
- 3 Congestion Management
- 4 Congestion Avoidance
- 5 Policing and shaping
- 6 Resource Reservation Protocol
- 7 Bibliography



Classification overview

To provide a preferential service to a type of traffic, it must be classified. Classification is done in 2 steps:

- ➊ Traffic must be identified. Identification methods:
 - Use of ACLs (Access Control Lists)
 - Definition of **route maps**
- ➋ Optionally may be marked.
 - If identified and not marked, classification is said to be on a per-hop basis. Not passed to the next router
 - When marked for network-wide use, IP Precedence bits are set

When marked, routers can use IP Precedence bits to:

- determine how WFQ and WRED methods manages the traffic
- use features such as **policy-based routing** or **committed access rate (CAR)**

Policy-based routing

PBR. How it works ?


Traffic flows can be configured, marked and routed accordingly.

- Incoming traffic is classified using ACLs or extended ACLs. (Based on IPs, port numbers, packet length, ...)
- IP Precedence bits are set according to classification
- Specific next-hop routers may be set



Policy-based routing

Example:



Ping from 10.0.0.2 to 10.0.1.2

```

Internet Protocol Version 4, Src: 10.0.0.2 (10.0.0.2), Dst: 10.0.1.2 (10.0.1.2)
  Version: 4
  Header length: 20 bytes
  + Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00: Not-ECT (Not ECN-C
  
```


```

Internet Protocol Version 4, Src: 10.0.0.2 (10.0.0.2), Dst: 10.0.1.2 (10.0.1.2)
  Version: 4
  Header length: 20 bytes
  + Differentiated Services Field: 0xe0 (DSCP 0x38: Class Selector 7; ECN: 0x00: Not-ECT (Not ECN-C
  
```



Policy-based routing

Example:



Ping from 10.0.0.3 to 10.0.1.2

```

Internet Protocol Version 4, Src: 10.0.0.3 (10.0.0.3), Dst: 10.0.1.2 (10.0.1.2)
  Version: 4
  Header length: 20 bytes
  Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00: Not-ECT (Nc
  
```

```

Internet Protocol Version 4, Src: 10.0.0.3 (10.0.0.3), Dst: 10.0.1.2 (10.0.1.2)
  Version: 4
  Header length: 20 bytes
  Differentiated Services Field: 0x20 (DSCP 0x08: Class Selector 1; ECN: 0x00: Not-E
  
```



Policy propagation via BGP

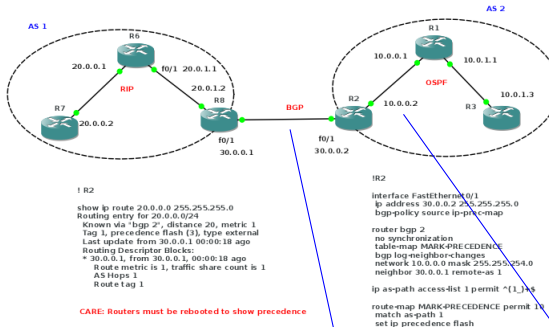
Configuration

Allows packet classification marking IP precedence based on BGP community

- 1 Indicate to the incoming i/f that bgp-policy IP precedence classification must be used
- 2 Define access list matching the required path
- 3 Define a route-map setting the IP precedence
- 4 Use the route-map defined in the BGP router instance

Policy propagation via BGP

Example:



Ping from 20.0.0.2 to 10.0.1.3

```
Internet Protocol Version 4, Src: 20.0.0.2 (20.0.0.2), Dst: 10.0.1.3 (10.0.1.3)
  Version: 4
  Header length: 20 bytes
  Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00: Not-ECT (No
    Total Length: 100
```

```
Internet Protocol Version 4, Src: 20.0.0.2 (20.0.0.2), Dst: 10.0.1.3 (10.0.1.3)
  Version: 4
  Header length: 20 bytes
  Differentiated Services Field: 0x60 (DSCP 0x18: Class Selector 3; ECN: 0x00: Not-ECT
    Total Length: 100
  Identification: 0x0000 (0)
```



Committed Access Rate (CAR)

CAR

CAR is a feature that implements classification and policing. Limits the input or output rate at an i/f.

Rate policies can be applied according to:

- All IP traffic
- IP precedence
- MAC address
- IP access list

Committed Access Rate (CAR)

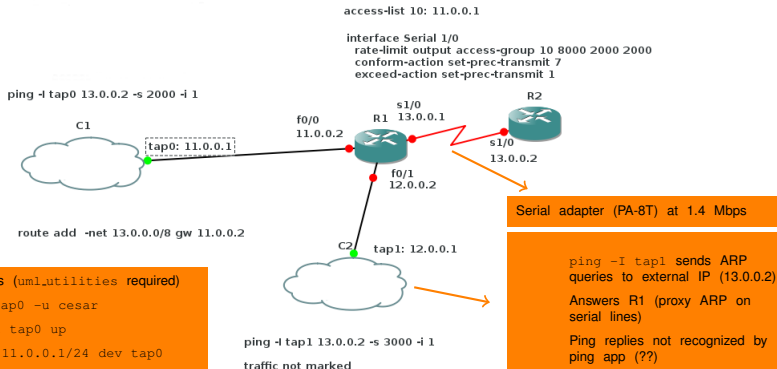
CAR configuration

Configuration is done in a interface:

- Set **rate-limit** for input or output traffic giving:
 - Average rate (in bps)
 - Normal burst size (in bytes)
 - Maximum burst size. Bursts between normal and maximum are considered exceeding with increasing probability
- Set the actions to be performed for conforming (**conform-action**) and exceeding (**exceed-action**) traffic. Actions can be:
 - Drop the packet
 - Transmit
 - Set precedence and transmit
 - Continue (evaluate the next rate-limit action)
 - Set the precedence and continue



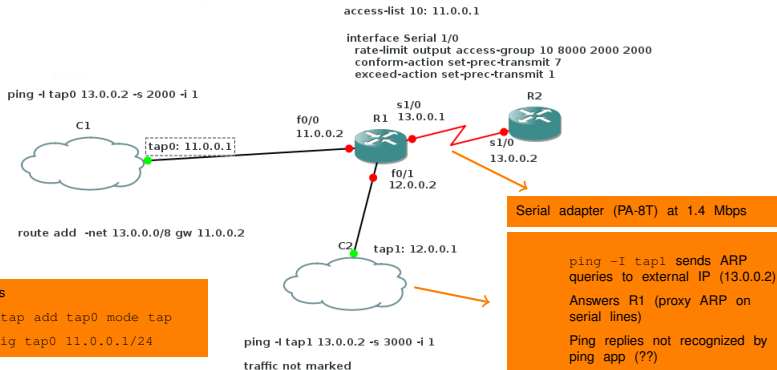
Committed Access Rate (CAR)



- Traffic from tap0. 2000 bytes every second. Rate 16 Kbps > 8 Kbps.
- Approx. half of the packets will be set to IP prec 1 outcoming s1/0
- No packets marked from tap1
- Check capture at serial line s1/0



Committed Access Rate (CAR)



CentOS: Tap interfaces

```

sudo ip tuntap add tap0 mode tap
sudo ifconfig tap0 11.0.0.1/24

```

- Traffic from tap0. 2000 bytes every second. Rate 16 Kbps > 8 Kbps.
- Approx. half of the packets will be set to IP prec 1 outcoming s1/0
- No packets marked from tap1
- Check capture at serial line s1/0



Contents

1 QoS Overview

2 Classification

3 Congestion Management

- Congestion management overview
- Flow-based WFQ
- Class-based WFQ (CBWFQ)
- Custom queueing (CQ)
- Priority queueing (PQ)
- Low Latency Queueing (LLQ)

4 Congestion Avoidance

5 Policing and shaping

6 Resource Reservation Protocol



Congestion management overview

Congestion management tasks

- 1 Creation of software queues
- 2 Assign packets to queues based on classification
- 3 Schedule packets in queues for transmission



Congestion management overview


Types of queues

- FIFO. No QoS
- WFQ. Default for slow speed i/fs. 2 types:
 - Flow-based WFQ. A flow is determined by IPs, protocol and port numbers of a connection. Configurable number of queues. 256 as default. No configuration required
 - Class-based WFQ (CBWFQ). Definition of **class-maps** based on access-lists. 1 queue per class. Up to 64 classes
- Custom Queueing (CQ). Allocates bandwidth for each class of traffic. 16 queues. Round robin scheduling (Weighted round robin, WRR)
- Priority Queueing (PQ). Packets from a priority are sent before all lower priorities. Ensures low latency requirements. 4 queues
- Low Latency Queueing (LLC). Adds PQ to flow-based WFQ or CBWFQ



Flow-based WFQ

Schema



Source: Cisco IOS Quality of Service Solutions Configuration Guide, Release 12.2

Flow-based WFQ

WFQ and IP precedence

Having $Nflows_j$ flows with a IP precedence value j and an assigned weight w_j , the assigned bandwidth ($1/r_i$) to a flow of precedence i is computed as:

$$\frac{1}{r_i} = \frac{w_i}{\sum_{j=0}^7 Nflows_j \cdot w_j}$$

As flows are added and ended, the allocated bandwidth changes continuously

Example (taking $w_i = i + 1$)

Having 5 flows; 2 with IP precedence value 0 (routine), and 3 with IP precedence 5 (critical), their assigned bandwidth results:

$$\frac{1}{r_0} = \frac{1}{2 \cdot 1 + 3 \cdot 6} = 1/20 = 0.05$$

$$\frac{1}{r_5} = \frac{6}{2 \cdot 1 + 3 \cdot 6} = 6/20 = 0.3$$



Flow-based WFQ

Configuring WFQ

- WFQ is configured as default control management for slow speed links (<2 Mbps)
- Command **fair-queue** run on i/f basis. 3 parameters:
 - 1 *congestive-discard-threshold*. Number of packets allowed in each queue. Default 64
 - 2 *dynamic-queues*. Number of WFQ queues. Power of 2. Default depends on i/f BW. 256 for links > 512 Kbps
 - 3 *reservable-queues*. Reserved to RSVP (Integrated Services) or CBWFQ (DiffServ), Default 0

Flow-based WFQ

Monitoring WFQ

```
R1#show queue Serial 1/0
```

```
Input queue: 0/75/0/0 (size/max/drops/flushes); Total output drops:
```

```
Queueing strategy: weighted fair
```

```
Output queue: 0/1000/64/0 (size/max total/threshold/drops)
```

```
Conversations 0/1/256 (active/max active/max total)
```

```
Reserved Conversations 0/0 (allocated/max allocated)
```

```
Available Bandwidth 1158 kilobits/sec
```

```
R1#show queue Serial 1/0
```

```
Input queue: 0/75/0/0 (size/max/drops/flushes); Total output drops:
```

```
Queueing strategy: weighted fair
```

```
Output queue: 63/1000/64/668800 (size/max total/threshold/drops)
```

```
Conversations 1/2/256 (active/max active/max total)
```

```
Reserved Conversations 0/0 (allocated/max allocated)
```

```
Available Bandwidth 1158 kilobits/sec
```

```
(depth/weight/total drops/no-buffer drops/interleaves) 63/4048/66880
```

```
Conversation 29, linktype: ip, length: 332
```

```
source: 12.0.0.1, destination: 13.0.0.2, id: 0x3FFA, ttl: 63, prot:
```



Flow-based WFQ

WFQ weights

Predefined weights (w_i) are the following:

IP prec.	Name	WFQ weight ($1/w_i$)
111	Network Control	4,048
110	Internet Control	4,626
101	Critical	5,397
100	Flash override	6,476
011	Flash	8,096
010	Immediate	10,794
001	Priority	16,192
000	Routine	32,384

Computed as

$$\frac{1}{w_i} = \frac{32,384}{\text{IP_Prec}_i + 1}$$


Weights only can be configured as DWFQ (Distributed WFQ) that runs on advanced processors



Flow-based WFQ

Example

```
packETHcli -i tap0 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap0-100.pcap -n 0 (1.1 Mbps)
```




```
packETHcli -i tap1 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap1-300.pcap -n 0 (2.7 Mbps)
```

- No further configuration required on i/f s1/0
- Packets at input i/fs are IP precedence marked after classification; f0/1 network(7), f0/0 routine(0)
- ping from tap0. Data length 92 bytes. ICMP header (8 bytes). IP header (20 bytes). Ethernet header (14 bytes). Total packet length: 134 bytes. **Data rate:** $134 \cdot 8 / 10^{-3} = 1.07 \text{ Mbps}$
- ping from tap1. Data length 292 bytes. Total packet length: 334 bytes. Data rate: 2.67 Mbps
- Find [packETH](#) [here](#) and [packETHcli](#) [here](#)

Flow-based WFQ

Example




- Assigned bandwidths:
 - From tap0 : $1/r_0 = \frac{1/32,384}{1/4,048 + 1/32,384} \simeq 1/9$
 - From tap1 : $1/r_1 = \frac{1/4,048}{1/4,048 + 1/32,384} \simeq 8/9$
- Capturing at s1/0 , taking 793 packets, we observe:
 - 189 packets from tap0 . 128 bytes each (IP packet). 24,192 bytes
 - 604 packets from tap1 . 198,112 bytes
 - Gives a ratio of $8.18 = 198,112 / 24,192$



Flow-based WFQ

Example

```
packETHcli -i tap0 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap0-100.pcap -n 0 (1.1 Mbps)
```




```
packETHcli -i tap1 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap1-300.pcap -n 0 (2.7 Mbps)
```

R1 configuration

```
interface FastEthernet0/0
 ip address 11.0.0.2 255.255.255.0
 ip policy route-map RM0
!
interface FastEthernet0/1
 ip address 12.0.0.2 255.255.255.0
 ip policy route-map RM1
```

Flow-based WFQ

Example



R1 configuration

```

access-list 1 permit any
!
route-map RM1 permit 10
  match ip address 1
  set ip precedence network
!
route-map RM0 permit 10
  match ip address 1
  set ip precedence routine


```



Flow-based WFQ

Example

```
packETHcli -i tap0 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap0-100.pcap -n 0 (1.1 Mbps)
```



```
packETHcli -i tap1 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap1-300.pcap -n 0 (2.7 Mbps)
```

Monitoring queues

```
R1#show queueing interface Serial 1/0
Interface Serial1/0 queueing strategy: fair
  Input queue: 0/75/0/0 (size/max/drops/flushes); Total output drops: 46388
  Queueing strategy: weighted fair
  Output queue: 64/1000/64/46388 (size/max total/threshold/drops)
    Conversations 2/3/256 (active/max active/max total)
    Reserved Conversations 0/0 (allocated/max allocated)
    Available Bandwidth 1158 kilobits/sec

(depth/weight/total drops/no-buffer drops/interleaves) 48/4048/22662/0/0
Conversation 29, linktype: ip, length: 332
source: 12.0.0.1, destination: 13.0.0.2, id: 0x3FFA, ttl: 63, prot: 1

(depth/weight/total drops/no-buffer drops/interleaves) 16/32384/23727/0/0
Conversation 28, linktype: ip, length: 132
source: 11.0.0.1, destination: 13.0.0.2, id: 0x3F9E, ttl: 63, prot: 1
```



Class-based WFQ (CBWFQ)

Characteristics

- Classes defined according to matching criteria, access-lists and input i/fs
- A single queue (from the WFQ) is reserved for each class
- Parameters to assign at each class-queue:
 - **Bandwidth.** In bps or a % of the total. A **max-reserved-bandwidth** is set as default to 75%
 - **Weight** for WFQ is automatically derived from the assigned bandwidth
 - **Queue limit**
- Packet drop. Once the queue limit is reached, **tail-drop** applies. WRED can be also configured
- Up to 64 class-queues per i/f



Class-based WFQ (CBWFQ)


Configuration steps

- Define **class-maps** specifying how traffic is classified
- Define **policy-maps** indicating what to do with defined classes
- Apply policies to i/fs

Class-based WFQ (CBWFQ)

Example

```
packETHcli 4 tap0 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap0-100.pcap -n 0 (1.1 Mbps)
```



```
packETHcli 4 tap1 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap1-300.pcap -n 0 (2.7 Mbps)
```

CBWFQ:

- 80% from s1/0 BW to IP traffic from tap0
- 20% from s1/0 BW to IP traffic from tap1

Requirements (IOS 12)


- 80% bandwidth of s1/0 reserved to IP traffic from tap0
- 20% bandwidth of s1/0 reserved to IP traffic from tap1
- Being so, max-reserved-bandwidth must be set to 100%



Class-based WFQ (CBWFQ)

Example

```
packETHcli 4 tap0 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap0-100.pcap -n 0 (1.1 Mbps)
```



```
packETHcli 4 tap1 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap1-300.pcap -n 0 (2.7 Mbps)
```

CBWFQ:

- 80% from s1/0 BW to IP traffic from tap0
- 20% from s1/0 BW to IP traffic from tap1

Requirements (IOS 15)


- max-reserved-bandwidth deprecated
- 80% bandwidth of s1/0 reserved to IP traffic from tap0
- 19% bandwidth of s1/0 reserved to IP traffic from tap1 (to avoid error: should be less than 100%)



Class-based WFQ (CBWFQ)

Example

```
packETHcli 4 tap0 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap0-100.pcap -n 0 (1.1 Mbps)
```



CBWFQ:

- 80% from s1/0 BW to IP traffic from tap0
- 20% from s1/0 BW to IP traffic from tap1

```
packETHcli 4 tap1 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap1-300.pcap -n 0 (2.7 Mbps)
```

R1 configuration


```
access-list 101 permit ip 11.0.0.0 0.0.0.255 any
access-list 102 permit ip 12.0.0.0 0.0.0.255 any
!
class-map match-all class1
  match access-group 101
class-map match-all class2
  match access-group 102
```

```
policy-map policy1
  class class1
    bandwidth percent 80
  class class2
    bandwidth percent 19
!
interface Serial1/0
  ip address 13.0.0.1 255.255.255.0
  %max-reserved-bandwidth 100
  service-policy output policy1
```

Class-based WFQ (CBWFQ)

Example

```
packETHcli 4 tap0 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap0-100.pcap -n 0 (1.1 Mbps)
```



CBWFQ:

- 80% from s1/0 BW to IP traffic from tap0
- 20% from s1/0 BW to IP traffic from tap1

```
packETHcli 4 tap1 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap1-300.pcap -n 0 (2.7 Mbps)
```

Run test

i/f	# Packets	IP length (bytes)	Traffic volume (bytes)
tap0	1,917	128	245,376
tap1	195	328	63,960

$$\text{Traffic ratio} = \frac{245,376}{63,960} = 3.83 \simeq \frac{80}{19}$$



Class-based WFQ (CBWFQ)

Queue monitoring

```
R1# show queueing interface Serial 1/0
Interface Serial1/0 queueing strategy: fair
Input queue: 0/75/0/0 (size/max/drops/flushes); Total output drops: 14010
Queueing strategy: Class-based queueing
Output queue: 129/1000/64/14010 (size/max total/threshold/drops)
Conversations 3/3/256 (active/max active/max total)
Reserved Conversations 2/2 (allocated/max allocated)
Available Bandwidth 1 kilobits/sec

(depth/weight/total drops/no-buffer drops/interleaves) 64/20/5634/0/0
Conversation 265, linktype: ip, length: 132
source: 11.0.0.1, destination: 13.0.0.2, id: 0x3F9E, ttl: 63, prot: 1

(depth/weight/total drops/no-buffer drops/interleaves) 64/78/8380/0/0
Conversation 266, linktype: ip, length: 332
source: 12.0.0.1, destination: 13.0.0.2, id: 0x3FFA, ttl: 63, prot: 1

(depth/weight/total drops/no-buffer drops/interleaves) 1/32384/0/0/0
Conversation 257, linktype: cdp, length: 333
```



Custom queueing (CQ)

Characteristics

- Up to 16 configurable queues
- Configurable parameters:
 - **limit**: max number of packet per queue (default 20)
 - **byte-count**: counts the number of bytes per queue served at each round. If limit is reached while transmitting a packet, the remaining packet is transmitted
- Queues served in a **round-robin** fashion
- Guaranteed bandwidth can be easily derived
- Packet **classification**: based on **protocol** type or **interfaces**
- A **default queue** can be assigned to non-matching traffic



Custom queueing (CQ)

Assigning the byte-count. An example

Assume we want to allocate 3 traffic flows as follows (IP lengths are known and supposed fixed):

Traffic	IP length (L) (bytes)	BW reserved (B) (%)
A	200	30
B	450	50
C	1,500	20

We proceed:


Traffic	B/L	Normalized B/L (N)	byte-count = N·L
A	0.150	11.2	2,240
B	0.111	8.3	3,735
C	0.013	1.0	1,500

Custom queueing (CQ)

Example

```
ping -s 16000 13.0.0.2 -c 1 4 tap0
```

Only 4 out of 16000/1500 packets get R2



```
ping -s 16000 13.0.0.2 -c 1 4 tap1
```

Only 7 out of 16000/1500 packets get R2

Queue monitoring

```
R1#show interfaces Serial 1/0
```

```
Serial1/0 is up, line protocol is up
```

```
....
```

```
Output queues: (queue #: size/max/drops)
```

```
0: 0/3/7 1: 0/6/3 2: 0/20/0 3: 0/20/0 4: 0/20/0
5: 0/20/0 6: 0/20/0 7: 0/20/0 8: 0/20/0 9: 0/20/0
10: 0/20/0 11: 0/20/0 12: 0/20/0 13: 0/20/0 14: 0/20/0
15: 0/20/0 16: 0/20/0
```




Custom queueing (CQ)

Example

```
ping -s 16000 13.0.0.2 -c 1 4 tap0
```

Only 4 out of 16000/1500 packets get R2



```
ping -s 16000 13.0.0.2 -c 1 4 tap1
```

Only 7 out of 16000/1500 packets get R2

Queue monitoring

```
R1# clear counters Serial 1/0
```

```
R1# show queueing interface Serial 1/0
```

```
Interface Serial1/0 queueing strategy: custom
```

```
Output queue utilization (queue/count)
```

```

0/4 1/7 2/0 3/0 4/0 5/0 6/0 7/0 8/0
9/0 10/0 11/0 12/0 13/0 14/0 15/0 16/0
  
```



Priority queueing (PQ)

Characteristics


- Up to 4 queues. **High, medium, normal** and **low**
- PQ gives **absolute** priority. Highest priority queues are first processed until being emptied
 - Some lowest priority packets may be never sent
 - Use traffic shaping or CAR to avoid previous issue on higher priorities
- Packets classified as usual
- Not classified packets ingress **normal** priority queue
- PQ adds extra processing. Not acceptable for high speed i/fs



Priority queueing (PQ)

Example. Video streaming

```
ffmpeg -re -i file.avi -vcodec copy -an -sdp_file s.sdp -f rtp rtp://13.0.0.1:5004
```



NAT configuration

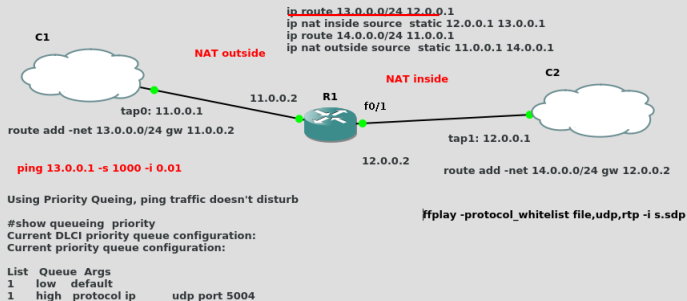
- Video streaming between 2 tap i/fs through R1
- tap1 seen as 13.0.0.1 from tap0
- tap0 seen as 14.0.0.1 from tap1



Priority queueing (PQ)

Example. Video streaming

```
ffmpeg -re -i file.avi -vcodec copy -an -sdp_file s.sdp -f rtp rtp://13.0.0.1:5004
```




Traffic configuration

- Video stream from tap0 to tap1 at 800 Kbps
- ping traffic would cause stream losses without PQ

Priority queueing (PQ)

Example. Video streaming

```
ffmpeg -re -i file.avi -vcodec copy -an -sdp_file s.sdp -f rtp rtp://13.0.0.1:5004
```



Using Priority Queing, ping traffic doesn't disturb

```
#show queueing priority
Current DLCI priority queue configuration:
Current priority queue configuration:
```

```
List Queue Args
1 low default
1 high protocol ip      udp port 5004
```

```
ffplay -protocol_whitelist file,udp,rtp -i s.sdp
```

R1 PQ configuration

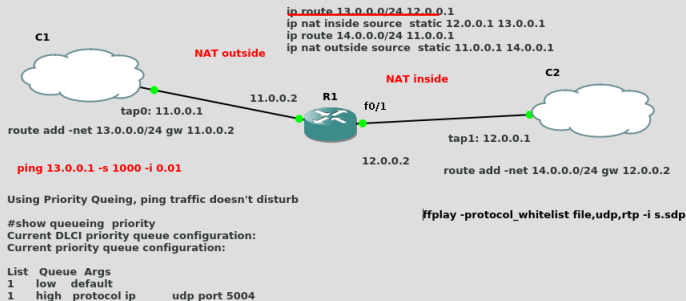
```
interface FastEthernet0/1
 ip address 12.0.0.2 255.255.255.0
 ip nat inside
 priority-group 1
 !
 priority-list 1 protocol ip high udp 5004
 priority-list 1 default low
```



Priority queueing (PQ)

Example. Video streaming

```
ffmpeg -re -i file.avi -vcodec copy -an -sdp_file s.sdp -f rtp rtp://13.0.0.1:5004
```



PQ monitoring

```
R1# show queueing interface FastEthernet 0/1
Interface FastEthernet0/1 queueing strategy: priority
```


```
Output queue utilization (queue/count)
high/2772 medium/0 normal/16254 low/862
```



Priority queueing (PQ)

Example. Video streaming

```
ffmpeg -re -i file.avi -vcodec copy -an -sdp_file s.sdp -f rtp rtp://13.0.0.1:5004
```



Using Priority Queing, ping traffic doesn't disturb

```
#show queueing priority
Current DLCI priority queue configuration:
Current priority queue configuration:
```

```
List Queue Args
1 low default
1 high protocol ip      udp port 5004
```

```
ffplay -protocol_whitelist file,udp,rtp -i s.sdp
```

PQ monitoring

```
R1# show queueing priority
Current DLCI priority queue configuration:
Current priority queue configuration:
```

```
List Queue Args
1 low default
1 high protocol ip      udp port 5004
```



Low Latency Queueing (LLQ)

Overview


- LLQ adds PQ to CBWFQ
- Useful for real-time applications such as audio calls. Reduce jitter in voice conversations
- Voice comms uses UDP, not suitable to WRED congestion avoidance
- Packets in PQ are dequeued before those in WFQ queues
- LLQ uses a single priority queue within the CBWFQ classes



Low Latency Queueing (LLQ)

Example

```
packETHcli -i tap0 -d 1000 -m 2 -f /home/cesar/GNS3/Captures/ping-tap0-100.pcap -n 0 (1.1 Mbps)
```



Low Latency Queueing (LLQ)

R1 configuration

```
access-list 101 permit ip 11.0.0.0 0.0.0.255 any
access-list 102 permit ip 12.0.0.0 0.0.0.255 any
access-list 103 permit ip 14.0.0.0 0.0.0.255 any

class-map match-all class1
  match access-group 101
class-map match-all class2
  match access-group 102
class-map match-all class3
  match access-group 103

policy-map policy1
  class class1
    bandwidth percent 60
  class class2
    bandwidth percent 20
  class class3
    priority percent 20 ← PQ
```



Low Latency Queueing (LLQ)

Queue monitoring

```

R1#show queueing interface Serial 1/0
Interface Serial1/0 queueing strategy: fair
  Input queue: 0/75/0/0 (size/max/drops/flushes); Total output drops: 215341
  Queueing strategy: Class-based queueing
  Output queue: 133/1000/64/215341 (size/max total/threshold/drops)
    Conversations 4/4/256 (active/max active/max total)
    Reserved Conversations 2/2 (allocated/max allocated)
    Available Bandwidth 2 kilobits/sec

    (depth/weight/total drops/no-buffer drops/interleaves) 5/0/106/0/0
    Conversation 264, linktype: ip, length: 1504
    source: 14.0.0.2, destination: 13.0.0.2, id: 0x0115, ttl: 254, prot: 1

    (depth/weight/total drops/no-buffer drops/interleaves) 64/26/103921/0/0
    Conversation 265, linktype: ip, length: 132
    source: 11.0.0.1, destination: 13.0.0.2, id: 0x3F9E, ttl: 63, prot: 1

    (depth/weight/total drops/no-buffer drops/interleaves) 64/78/111317/0/0
    Conversation 266, linktype: ip, length: 332
    source: 12.0.0.1, destination: 13.0.0.2, id: 0x3FFA, ttl: 63, prot: 1

    (depth/weight/total drops/no-buffer drops/interleaves) 3/32384/0/0/0
    Conversation 33, linktype: ip, length: 104
    source: 15.0.0.2, destination: 13.0.0.2, id: 0x001F, ttl: 254, prot: 1

```

PQ weight



Low Latency Queueing (LLQ)

ping from R3

```
R3# ping 13.0.0.2 size 100 repeat 10
```

```
Type escape sequence to abort.
```

```
Sending 10, 100-byte ICMP Echos to 13.0.0.2, timeout is 2 seconds:
```

```
!!!!!!!!!!!!
```

```
Success rate is 100 percent (10/10), round-trip min/avg/max = 52/60/68
```

ping from R4

```
R4 #ping 13.0.0.2 size 100 repeat 10
```

```
Type escape sequence to abort.
```

```
Sending 10, 100-byte ICMP Echos to 13.0.0.2, timeout is 2 seconds:
```

```
.....
```

```
Success rate is 0 percent (0/10)
```



Contents

- 1 QoS Overview
- 2 Classification
- 3 Congestion Management
- 4 Congestion Avoidance**
 - Congestion avoidance overview
 - WRED
- 5 Policing and shaping
- 6 Resource Reservation Protocol
- 7 Bibliography



Congestion avoidance overview


Congestion management tasks

- **RED** is used to prevent congestion
- **Tail-drop** as **default** if no RED configured
- CISCO implements a **weighted** version of RED (**WRED**), combining RED and IP Precedence. Weighted can be disabled, turning into a simple RED mechanism
- WRED additional features:
 - **Flow-based WRED**. More fairness to all flows
 - **Diffserv WRED**. Drop probabilities based on differentiated service code points (DCSP)

Congestion avoidance overview

RED fundamentals

- See chapter 1 (TCP congestion) to revisit RED mechanics
- Only effective in TCP flows
- Drops cause TCP not increasing advertise windows. Too much drops can put TCP into slow start
- Parameter names in CISCO configuration:



Computing the average queue length:

$$\text{AvgLen} = (1 - 2^{-n}) \cdot \text{AvgLen} + 2^{-n} \cdot \text{SampleLen}$$

n: exponential-weighting-constant



WRED


WRED basics

- A different probability profile applied to each IP precedence
- To turn WRED into RED, put the same values to all IP precedences

WRED default values

Exponential weighting constant (n): 9


Class	Min. Thresh.
0	20
1	22
2	24
3	26
4	28
5	31
6	33
7	35
RSVP	37



WRED

Example

tcpspray in ndisc6 package(Fedora)




R1 configuration

```
interface Serial1/0
ip address 13.0.0.1 255.255.255.0
random-detect
```


WRED

Example



R1 configuration

```

interface FastEthernet0/0
 ip address 11.0.0.2 255.255.255.0
 ip policy route-map RM0

```

```

access-list 100 permit tcp 11.0.0.0 0.0.0.255 14.0.0.0 0.0.0.255
access-list 110 permit tcp 11.0.0.0 0.0.0.255 15.0.0.0 0.0.0.255

```

```

route-map RM0 permit 10
 match ip address 100
 set ip precedence priority
!
route-map RM0 permit 20
 match ip address 101
 set ip precedence network

```


WRED

Example

```
tcpspray 14.0.0.2 -n 1000
tcpspray 15.0.0.2 -n 1000
```

C1

```
route add -net 14.0.0.0/24 gw 11.0.0.2
route add -net 15.0.0.0/24 gw 11.0.0.2
```



RED monitoring


```
R1#show queueing interface Serial 1/0
Interface Serial1/0 queueing strategy: random early detection (WRED)
Random-detect not active on the dialer
Exp-weight-constant: 9 (1/512)
Mean queue depth: 0
```

class	Random drop pkts/bytes	Tail drop pkts/bytes	Minimum thresh	Maximum thresh	Mark prob
0	0/0	0/0	20	40	1/10
1	0/0	0/0	22	40	1/10
2	0/0	0/0	24	40	1/10
3	0/0	0/0	26	40	1/10
4	0/0	0/0	28	40	1/10
5	0/0	0/0	31	40	1/10
6	0/0	0/0	33	40	1/10
7	0/0	0/0	35	40	1/10

tcpspray doesn't push
to the limit

WRED

Example



Modifying WRED parameters

```


ip address 13.0.0.1 255.255.255.0
random-detect
random-detect precedence 7 23 40 10
  
```

IP prec min-thresh max-thresh mark-prob denominator

WRED

Example 2

```
scp arxiu cesar@22.0.0.1/tmp
scp arxiu cesar@23.0.0.1/tmp
```



RED monitoring

```
R1#show queueing random-detect
Current random-detect configuration:
FastEthernet0/1
  Queueing strategy: random early detection (WRED)
  Random-detect not active on the dialer
  Exp-weight-constant: 9 (1/512)
  Mean queue depth: 0
```


class	Random drop pkts/bytes	Tail drop pkts/bytes	Minimum thresh	Maximum thresh	Mark prob
0	0/0	0/0	20	40	1/10
1	128/193792	0/0	22	40	1/10
2	0/0	0/0	24	40	1/10
3	0/0	0/0	26	40	1/10
4	0/0	0/0	28	40	1/10
5	0/0	0/0	31	40	1/10
6	28/42392	0/0	33	40	1/10
7	0/0	0/0	35	40	1/10
rsdp	0/0	0/0	37	40	1/10



WRED

Example 2

```
scp arxiu cesar@22.0.0.1:/tmp
scp arxiu cesar@23.0.0.1:/tmp
```



RED throughput

```
cesar@pdm:/
%
% scp tlou8.mkv cesar@22.0.0.1:/tmp/
cesar@22.0.0.1's password:
tlou8.mkv                               1%   39MB 682.8KB/s 1:24:12 ETA

cesar@pdm:/
%
% scp tlou8.mkv cesar@23.0.0.1:/tmp/
cesar@23.0.0.1's password:
tlou8.mkv                               2%   80MB 1.3MB/s 41:20 ETA
```



Contents

- 1 QoS Overview
- 2 Classification
- 3 Congestion Management
- 4 Congestion Avoidance
- 5 Policing and shaping**
 - Policing and shaping overview
 - Token bucket
 - Traffic policing
 - Traffic shaping
- 6 Resource Reservation Protocol
- 7 Bibliography



Policing and shaping overview

Policing and shaping are **traffic regulation** mechanisms

- **Policing**: non-compliant traffic is discarded. (ex. CAR policy seen before)
- **Shaping**: non-compliant traffic is shaped and transmitted

How compliance is determined ? : **Token bucket**



Token bucket

Definition

Token bucket is a formal definition of **data transfer rate**. 3 components:

- **Mean rate** (r): amount of data to be transferred per unit time **on average**. Also called *Committed Information Rate* (CIR)
- **Burst size** (b): amount of data that **can be** transferred in a given **time interval**
- **Time interval** (t): Burst size expressed in time units. Derived from mean rate and burst size

The following relation holds: $r = \frac{b}{t}$

Over any integral part of t , transmit rate must not exceed r . Inside t , may be arbitrarily fast.



Token bucket


Algorithm

- A token is added to the bucket every $\frac{1}{r}$ seconds
- The bucket size is b . Tokens arriving when the bucket is full are discarded
- An arriving packet of size d is determined as:
 - **conformant**, if d is smaller than the number of tokens in the bucket
 - **non-conformant**, otherwise
- Conformant packets are transmitted. Bucket is decremented in d
- Non-conformant packets are discarded (policing) or delayed until enough tokens (shaping)



Token bucket

Algorithm



Token bucket

Token bucket with 2 buckets

- Exceeding bursts may be allowed
- An **exceed bucket** is added to the already existing **conform bucket**
- Overflowing tokens from conform bucket drops into the exceed bucket
- Tokens can be borrowed from exceed bucket if conform bucket is not enough
- In this case, 3 actions must be observed:
 - **Conformant**
 - **Exceeded**
 - **Violated**
- Conform bucket size (b_c). Exceed bucket size (b_e)
- $b_e = 0$ is a token bucket with 1 bucket



Token bucket


Algorithm with two buckets

- A token is added to the conform bucket every $\frac{1}{r}$ seconds
- Tokens arriving when the conform bucket is full fills into the exceed bucket
- An arriving packet of size d is determined as:
 - **conformant**, if d is smaller than the number of tokens in the conform bucket
 - **exceed**, if d is greater than the number of tokens in the conform bucket and smaller than the number of tokens in the conform and exceed bucket
 - **violated**, otherwise
- Conformant packets are transmitted. Conform bucket is decremented in d
- Exceeded packets are treated according to exceeding policy. Conform bucket is emptied and exceed bucket is decremented in $(d - b)$
- Violated packets are treated according to violating policy



Token bucket

Algorithm with 2 buckets



Token bucket

A numerical example

Tokens rate = 1. $b_c = 4$. $b_e = 6$

time	Packet Length	Conform bucket	Exceed bucket	Action
0	-	4	6	-
1	2	3	6	conform
2	-	4	6	-
3	5	1	5	exceed
4	-	2	5	-
5	5	1	2	exceed
6	4	2	2	violated
7	-	3	2	-
8	-	4	2	-
9	-	4	3	-

Traffic policing

Overview

- Traffic policing allows control of maximum **incoming or leaving** rate using token bucket
- Traffic can be partitioned into several classes
- Several actions on conforming, exceeding and violating traffic:
 - Drop
 - Transmit
 - Set IP precedence and transmit
 - Set DSCP value and transmit



Traffic policing

Configuration steps

- Configure a **class map**
- Configure a **police map**
- Configure token bucket parameters inside a policy map:
 - Average rate (in bps or as a fraction of the bandwidth)
 - Conformant bucket size (b_c) (in bytes)
 - Excess burst parameter ($b_c + b_e$) (in bytes). If excess burst parameter equals b_c , then $b_e = 0$
 - Conform, exceed and violate actions



Traffic policing

Configuration example

CAR example ([here](#)) can be also configured as follows:

```
access-list 1 permit 11.0.0.1
class-map match-all CMA
  match access-group 1

policy-map PMa
  class CMA
    police 8000 2000 2000 conform-action set-prec-transmit 7
    exceed-action set-prec-transmit 1

interface Serial1/0
  ip address 13.0.0.1 255.255.255.0
  service-policy output PMa
```



Traffic shaping

Overview


- Traffic shaping allows modify the **leaving** traffic profile to commit a given rate using token bucket
- Being so, we ensure traffic conforms certain policies
- As a result, traffic may suffers delays
- Shaping may be done based on ACLs or traffic classes



Traffic shaping

Shaping example. Video streaming

```
ffmpeg -re -i file.avi -vcodec copy -an -sdp_file s.sdp -f rtp rtp://13.0.0.1:5004
```




- Video stream at a 800 Kbps rate
- Shaping configured at i/f f0/1 at R1. 800 Kbps mean rate. No exceed bucket

Traffic shaping

Shaping example. Video streaming

```
ffmpeg -re -i file.avi -vcodec copy -an -sdp_file s.sdp -f rtp rtp://13.0.0.1:5004
```



R1 config

```

interface FastEthernet0/1
 ip address 12.0.0.2 255.255.255.0
 ip nat inside
 traffic-shape rate 800000 100000 100000

```

Average rate (bps)


b_c (bits)

$b_c + b_e$ (bits)



Traffic shaping

Traffic at R1 f0/1



Traffic shaping

Traffic shaping monitoring

```
R1#show traffic-shape
```

```
Interface Fa0/1
```

VC	Access List	Target Rate	Byte Limit	Sustain bits/int	Excess bits/int	Interval (ms)	Increment (bytes)
-		800000	25000	100000	100000	125	12500

r : average rate (bps)

b_c (bits)

$b_c + b_e$ (bits)

t : time interval (ms)

$$r = \frac{b_c}{t}$$

Traffic shaping

Traffic shaping monitoring

R1# show traffic-shape statistics

I/F	Acc. Queue	Packets	Bytes	Packets	Bytes	Shaping
	List Depth			Delayed	Delayed	Active
Fa0/1	0	3246	4342043	592	805842	no

R1# show traffic-shape statistics

I/F	Acc. Queue	Packets	Bytes	Packets	Bytes	Shaping
	List Depth			Delayed	Delayed	Active
Fa0/1	0	3268	4372183	592	805842	no

R1# show traffic-shape statistics

I/F	Acc. Queue	Packets	Bytes	Packets	Bytes	Shaping
	List Depth			Delayed	Delayed	Active
Fa0/1	0	3308	4426983	595	809952	yes



Contents

- 1 QoS Overview
- 2 Classification
- 3 Congestion Management
- 4 Congestion Avoidance
- 5 Policing and shaping
- 6 Resource Reservation Protocol**
 - Overview
 - The protocol
 - Features
 - Configuration and monitoring

- 7 Bibliography



Overview

What is it ?

- RSVP is an implementation of **Integrated Services**
- **Signaling protocol**. By itself doesn't provide QoS
- Clients use RSVP to apply for QoS for a **session**
- A session consists of:
 - Destination address (unicast or multicast)
 - IP protocol
 - Destination port
- Layer 4 protocol on top IP
- Routers must implement QoS through WFQ, WRED, LLQ, ...



Overview

how does it work ? (I)

- RSVP reserves in only one direction (origin → destination)
- A route between **origin** and **destination** is established by routing protocols
- All the traversed routers must be informed about reservations
- Basic steps:
 - 1 **PATH** message initiated by origin to destination. This message includes:
 - Session parameters
 - QoS requirements (required rate, burst, delay, ...)
 - Nodes IP traversed
 - 2 **RESV** message from destination to origin following the same route (reverse) that PATH.
 - At each node (router), resource availability is determined
 - If enough resources, RESV message is forwarded to next node
 - If reservation is declined (not enough resources), an error message is sent to origin




Overview

how does it work ? (II)

- **Session maintenance.** PATH and RESV messages are periodically refreshed. 30" by default
- **Confirmation.** At each node, after a RESV, a CONFIRM message is sent to destination
- **Finishing reservations.** Can be finished by origin, destination or intermediate nodes
 - **PathTear** messages sent in PATH direction
 - **ResvTear** messages sent in RESV direction
- Tear messages **free** resources




The protocol



PATH message. Asks for a reservation: 1,000 Kbps (average rate) and 100 KBytes (Burst size). Protocol 1 (ICMP), any port (0)

```
Internet Protocol Version 4, Src: 11.0.0.1 (11.0.0.1), Dst: 13.0.0.1 (13.0.0.1)
Resource Reservation Protocol (RSVP): PATH Message. SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
+ RSVP Header. PATH Message.
+ SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
+ HOP: IPv4, 11.0.0.1
+ TIME VALUES: 30000 ms ← 30" refresh time
+ SENDER TEMPLATE: IPv4, Sender 11.0.0.1, Port 0.
+ SENDER TSPEC: IntServ, Token Bucket, 125000 bytes/sec. ← 1,000 Kbps
+ ADSPEC
```


The protocol



PATH forwarded to next node. Hop field changed

```
Internet Protocol Version 4, Src: 11.0.0.1 (11.0.0.1), Dst: 13.0.0.1 (13.0.0.1)
Resource ReserVation Protocol (RSVP): PATH Message. SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
+ RSVP Header. PATH Message.
+ SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
+ HOP: IPv4, 172.16.0.1
+ TIME VALUES: 30000 ms
+ SENDER TEMPLATE: IPv4, Sender 11.0.0.1, Port 0.
+ SENDER TSPEC: IntServ, Token Bucket, 125000 bytes/sec.
+ ADSPEC
```


The protocol



PATH forwarded to destination node

- Internet Protocol Version 4, Src: 11.0.0.1, Dst: 13.0.0.1
- Resource Reservation Protocol (RSVP): **PATH** Message. SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
 - RSVP Header. PATH Message.
 - SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
 - HOP: IPv4, 13.0.0.2**
 - TIME VALUES: 30000 ms
 - SENDER TEMPLATE: IPv4, Sender 11.0.0.1, Port 0.
 - SENDER TSPEC: IntServ, Token Bucket, 125000 bytes/sec.
 - ADSPEC


The protocol



RESV sent to **next hop**. Only 500 Kbps reserved

```
Internet Protocol Version 4, Src: 13.0.0.1 (13.0.0.1), Dst: 13.0.0.2 (13.0.0.2)
Resource Reservation Protocol (RSVP): RESV Message. SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
+ RSVP Header. RESV Message.
+ SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
+ HOP: IPv4, 13.0.0.1
+ TIME VALUES: 30000 ms
+ STYLE: Wildcard Filter (17)
+ FLOWSPEC: Controlled Load: Token Bucket, 62500 bytes/sec.
```

The protocol



CONFIRM to **destination** if enough resources on R4. RESV sent to **next hop**

```


} Internet Protocol Version 4, Src: 13.0.0.2 (13.0.0.2), Dst: 13.0.0.1 (13.0.0.1)
} Resource Reservation Protocol (RSVP): CONFIRM Message. SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
  + RSVP Header. CONFIRM Message.
  + SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
  + ERROR: IPv4, Error code: Confirmation, Value: 0, Error Node: 12.0.0.1
  + CONFIRM: Receiver 13.0.0.1
  + STYLE: Wildcard Filter (17)
  + FLOWSPEC: Controlled Load: Token Bucket, 62500 bytes/sec.

} Internet Protocol Version 4, Src: 172.16.0.2 (172.16.0.2), Dst: 172.16.0.1 (172.16.0.1)
} Resource Reservation Protocol (RSVP): RESV Message. SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
  + RSVP Header. RESV Message.
  + SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
  + HOP: IPv4, 172.16.0.2
  + TIME VALUES: 30000 ms
  + SCOPE

```



The protocol




CONFIRM to **destination** if enough resources on R2. RESV sent to **next hop**

- Internet Protocol Version 4, Src: 172.16.0.1 (172.16.0.1), Dst: 13.0.0.1 (13.0.0.1)
- Resource Reservation Protocol (RSVP): CONFIRM Message. SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
 - ⊕ RSVP Header. CONFIRM Message.
 - ⊕ SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
 - ⊕ ERROR: IPv4, Error code: Confirmation, Value: 0, Error Node: 12.0.0.1
 - ⊕ CONFIRM: Receiver 13.0.0.1
 - ⊕ STYLE: Wildcard Filter (17)
 - ⊕ FLOWSPEC: Controlled Load: Token Bucket, 62500 bytes/sec.
- Internet Protocol Version 4, Src: 11.0.0.2 (11.0.0.2), Dst: 11.0.0.1 (11.0.0.1)
- Resource Reservation Protocol (RSVP): RESV Message. SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
 - ⊕ RSVP Header. RESV Message.
 - ⊕ SESSION: IPv4, Destination 13.0.0.1, Protocol 1, Port 0.
 - ⊕ HOP: IPv4, 11.0.0.2
 - ⊕ TIME VALUES: 30000 ms
 - ⊕ STYLE: Wildcard Filter (17)
 - ⊕ FLOWSPEC: Controlled Load: Token Bucket, 62500 bytes/sec.



The protocol



Reservation at R2

To	From	Pro	DPort	Sport	Next Hop	I/F	Fi	Serv	BPS
13.0.0.1	11.0.0.1	1	0	0	172.16.0.2	Se1/0	WF	LOAD	500K


Reservation at R4

To	From	Pro	DPort	Sport	Next Hop	I/F	Fi	Serv	BPS
13.0.0.1	11.0.0.1	1	0	0	13.0.0.1	Fa0/0	WF	LOAD	500K

Reserved resources at output interfaces from origin to destination



The protocol



After a second reservation accepted from R6 results:

Reservation at R2

To	From	Pro	DPort	Sport	Next Hop	I/F	Fi	Serv	BPS
13.0.0.1	11.0.0.1	1	0	0	172.16.0.2	Sel1/0	WF	LOAD	500K
13.0.0.1	12.0.0.1	1	0	0	172.16.0.2	Sel1/0	WF	LOAD	500K

Reservation at R4

To	From	Pro	DPort	Sport	Next Hop	I/F	Fi	Serv	BPS
13.0.0.1	11.0.0.1	1	0	0	13.0.0.1	Fa0/0	WF	LOAD	500K
13.0.0.1	12.0.0.1	1	0	0	13.0.0.1	Fa0/0	WF	LOAD	500K



Features

Integrated services

- 2 type of services can be reserved:
 - Guaranteed-rate
 - Controlled-load
- **Guaranteed-rate.** Offered service as an unloaded network according to bandwidth requirements. Delay tolerant services. CISCO implements it using WFQ with weights proportional to bandwidth
- **Controlled-load.** Delivers assured bandwidth with constant delay. Implemented with WRED (not confirmed by experimentation)
- Both types of service may use LLQ. Reservations with rate and burst size below some threshold are considered priority and put into **priority queue** (Assigned weight is 0)



Features

Reservation styles

- A reservation belongs to a **class** and a **scope**
- Two classes:
 - 1 **Shared**. A single reservation is made for **multiple upstream senders**
 - 2 **Distinct**. A reservation established for each sender
- Two scopes:
 - 1 **Explicit**. The reservation is defined by a explicit **list of senders**
 - 2 **Wildcard**. Some wildcard (0) used to define multiple senders
- A **sender** consists of an origin IP and origin port



Features

Reservation styles

Such a combination of classes and scopes leads to **three** reservation styles

Scope	Classes	
	Distinct	Shared
Explicit	fixed-filter (FF)	shared-explicit (SE)
Wildcard	-	Wildcard-filter (WF)

Features

Reservation styles

Such a combination of classes and scopes leads to **three** reservation styles

Scope	Classes	
	Distinct	Shared
Explicit	fixed-filter (FF)	shared-explicit (SE)
Wildcard	-	Wildcard-filter (WF)

FF:

- Reservation not shared by any other senders
- If another **receiver** is added for the same sender, reservations are **merged**
- Example: video broadcast



Features

Reservation styles

Such a combination of classes and scopes leads to **three** reservation styles

Scope	Classes	
	Distinct	Shared
Explicit	fixed-filter (FF)	shared-explicit (SE)
Wildcard	-	Wildcard-filter (WF)

SE:

- Reservation shared by other senders
- Senders explicitly specified by the receiver

Features

Reservation styles

Such a combination of classes and scopes leads to **three** reservation styles

Scope	Classes	
	Distinct	Shared
Explicit	fixed-filter (FF)	shared-explicit (SE)
Wildcard	-	Wildcard-filter (WF)

WF:

- Reservation shared by other senders
- Senders specified by a wildcard
- WF and SE reservations useful for audio conference multicast.
No more than one link active at the same time




Configuration and monitoring

Example 1

```
while true;
do packetcli -i tap0 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/udp-size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 185;
sleep 1; done
```

```
ip rsvp reservation-host 13.0.0.1 11.0.0.1 udp 5004 0 wf load 1000 10
```

```
saved at ram as: ip rsvp reservation-host 13.0.0.1 0.0.0.0 udp 5004 0 wf load 1000 10
```



```
ip rsvp sender 13.0.0.1 11.0.0.1 udp 5004 0 11.0.0.1 FastEthernet0/0 1000 10
```

```
packetcli -i tap1 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/ping-1000-tap1.pcap -n 0
```

Enabling RSVP at interfaces (R2)

```
interface FastEthernet0/0
ip address 11.0.0.2 255.255.255.0
ip rsvp bandwidth 1200 1200
!
interface Serial1/0
ip address 172.16.0.1 255.255.255.0
ip rsvp bandwidth 1150 1150
```

Maximum amount of
reservable BW per flow (Kbps)

Default: previous value

Maximum amount of reservable BW per i/f.

75% of i/f rate as default
and as a maximum (Kbps)




Configuration and monitoring

Example 1

```
while true;
do packetcli 4 tap0 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/udp-size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 185;
sleep 1; done
```

```
ip rsvp reservation-host 13.0.0.1 11.0.0.1 udp 5004 0 wf load 1000 10
saved at ram as: ip rsvp reservation-host 13.0.0.1 0.0.0.0 udp 5004 0 wf load 1000 10
```



```
packetcli 4 tap1 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/ping-1000-tap1.pcap -n 0
```

Enabling RSVP at interfaces (R4)

```
interface FastEthernet0/0
ip address 13.0.0.2 255.255.255.0
ip rsvp bandwidth 1200 1200
!
interface Serial1/0
ip address 172.16.0.2 255.255.255.0
ip rsvp bandwidth 1150 1150
```

Enabling RSVP at interfaces (R3)

```
interface FastEthernet0/0
ip address 13.0.0.1 255.255.255.0
ip rsvp bandwidth 1200 1200
```




Configuration and monitoring

Example 1

```
while true;
do packetcli 4 tap0 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/udp-size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 185;
sleep 1; done
```

```
ip rsvp reservation-host 13.0.0.1 11.0.0.1 udp 5004 0 wf load 1000 10
saved at ram as: ip rsvp reservation-host 13.0.0.1 0.0.0.0 udp 5004 0 wf load 1000 10
```



RSVP PATH proxy (R2)

For clients without RSVP capabilities

```
ip rsvp sender 13.0.0.1 11.0.0.1 UDP 5004 0 11.0.0.1 FastEthernet0/0 1000 10
```

Prev. Hop

BW asked (kbps)

Max burst asked (KB)




Configuration and monitoring

Example 1

```
while true;
do packetcli 4 tap0 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/udp-size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 185;
sleep 1; done
```

```
ip rsvp reservation-host 13.0.0.1 11.0.0.1 udp 5004 0 wf load 1000 10
saved at ram as: ip rsvp reservation-host 13.0.0.1 0.0.0.0 udp 5004 0 wf load 1000 10
```



RSVP CONFIRM (R3)

```
ip rsvp reservation-host 13.0.0.1 11.0.0.1 UDP 5004 0 WF LOAD 1000 10
```

Reserv. style


BW reserved (kbps)

Max burst reserved (KB)



Configuration and monitoring

Example 1



Generate reserved traffic (tap0)

Average Rate: $\frac{1356 \cdot 8}{6000 \cdot 10^{-6}} = 1.8 \text{ Mbps}$

Burst size: $185 \cdot 1356 = 250 \text{ KB}$


Burst duration: $185 \cdot 6 \text{ ms} = 1.11 \text{ s}$

Configuration and monitoring

Example 1

```
while true;
do packetcli 4 tap0 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/udp-size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 185;
sleep 1; done
```

```
ip rsvp reservation-host 13.0.0.1 11.0.0.1 udp 5004 0 wf load 1000 10
saved at ram as: ip rsvp reservation-host 13.0.0.1 0.0.0.0 udp 5004 0 wf load 1000 10
```



```
ip rsvp sender 13.0.0.1 11.0.0.1 udp 5004 0 11.0.0.1 FastEthernet0/0 1000 10
```

```
packetcli 4 tap1 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/ping-1000-tap1.pcap -n 0
```


Generate non reserved traffic (tap1)
 ≈ 1.3 Mbps

Configuration and monitoring

Example 1

```
while true;
do packetcli & tap0 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/udp-size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 185;
sleep 1; done
```

```
ip rsvp reservation-host 13.0.0.1 11.0.0.1 udp 5004 0 wf load 1000 10
saved at ram as: ip rsvp reservation-host 13.0.0.1 0.0.0.0 udp 5004 0 wf load 1000 10
```



```
ip rsvp sender 13.0.0.1 11.0.0.1 udp 5004 0 11.0.0.1 FastEthernet0/0 1000 10
```

```
packetcli & tap1 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/ping-1000-tap1.pcap -n 0
```

Monitoring reservations (R2 and R4)

```
R2# show ip rsvp reservation
```

To	From	Pro	DPort	Sport	Next Hop	I/F	Fi	Serv	BPS
13.0.0.1	0.0.0.0	UDP	5004	0	172.16.0.2	Se1/0	WF	LOAD	1M

```
R2# show ip rsvp installation
```

```
RSVP: FastEthernet0/0 has no installed reservations
```

```
RSVP: Serial1/0
```

BPS	To	From	Protoc	DPort	Sport	Weight	Conversation
1M	13.0.0.1	0.0.0.0	UDP	5004	0	6	265

Configuration and monitoring

Example 1. Monitoring queues (R2)

```
R2#show queueing interface Serial 1/0
Interface Serial1/0 queueing strategy: fair
  Input queue: 0/75/0/0 (size/max/drops/flushes); Total output drops: 18256
  Queueing strategy: weighted fair
  Output queue: 106/1000/64/18256 (size/max total/threshold/drops)
    Conversations 3/4/256 (active/max active/max total)
    Reserved Conversations 1/1 (allocated/max allocated)
    Available Bandwidth 158 kilobits/sec
```

```
(depth/weight/total drops/no-buffer drops/interleaves) 9/6/0/0/0
Conversation 265, linktype: ip, length: 1360
source: 11.0.0.1, destination: 13.0.0.1, id: 0xEF7D, ttl: 63,
TOS: 0 prot: 17, source port 48823, destination port 5004
```

```
(depth/weight/total drops/no-buffer drops/interleaves) 42/32384/4304/0/0
Conversation 63, linktype: ip, length: 1360
source: 11.0.0.1, destination: 13.0.0.1, id: 0xEF7D, ttl: 63,
TOS: 0 prot: 17, source port 48823, destination port 5004
```

```
(depth/weight/total drops/no-buffer drops/interleaves) 55/32384/7187/0/0
Conversation 28, linktype: ip, length: 1032
source: 12.0.0.1, destination: 13.0.0.1, id: 0x39A8, ttl: 61, prot: 1
```

Reserved flow. Weight 6 in WFQ

Configuration and monitoring

Example 1. Monitoring queues (R2)

```
R2#show queueing interface Serial 1/0
Interface Serial1/0 queueing strategy: fair
  Input queue: 0/75/0/0 (size/max/drops/flushes); Total output drops: 18256
  Queueing strategy: weighted fair
  Output queue: 106/1000/64/18256 (size/max total/threshold/drops)
    Conversations 3/4/256 (active/max active/max total)
    Reserved Conversations 1/1 (allocated/max allocated)
    Available Bandwidth 158 kilobits/sec

    (depth/weight/total drops/no-buffer drops/interleaves) 9/6/0/0/0
    Conversation 265, linktype: ip, length: 1360
    source: 11.0.0.1, destination: 13.0.0.1, id: 0xEF7D, ttl: 63,
    TOS: 0 prot: 17, source port 48823, destination port 5004

    (depth/weight/total drops/no-buffer drops/interleaves) 42/32384/4304/0/0
    Conversation 63, linktype: ip, length: 1360
    source: 11.0.0.1, destination: 13.0.0.1, id: 0xEF7D, ttl: 63,
    TOS: 0 prot: 17, source port 48823, destination port 5004

    (depth/weight/total drops/no-buffer drops/interleaves) 55/32384/7187/0/0
    Conversation 28, linktype: ip, length: 1032
    source: 12.0.0.1, destination: 13.0.0.1, id: 0x39A8, ttl: 61, prot: 1
```

Not reserved flow. Weight 32,384 in WFQ. Best-effort

Configuration and monitoring

Example 1. Monitoring queues (R2)

```
R2#show queueing interface Serial 1/0
Interface Serial1/0 queueing strategy: fair
  Input queue: 0/75/0/0 (size/max/drops/flushes); Total output drops: 18256
  Queueing strategy: weighted fair
  Output queue: 106/1000/64/18256 (size/max total/threshold/drops)
    Conversations 3/4/256 (active/max active/max total)
    Reserved Conversations 1/1 (allocated/max allocated)
    Available Bandwidth 158 kilobits/sec

    (depth/weight/total drops/no-buffer drops/interleaves) 9/6/0/0/0
    Conversation 265, linktype: ip, length: 1360
    source: 11.0.0.1, destination: 13.0.0.1, id: 0xEF7D, ttl: 63,
    TOS: 0 prot: 17, source port 48823, destination port 5004


    (depth/weight/total drops/no-buffer drops/interleaves) 42/32384/4304/0/0
    Conversation 63, linktype: ip, length: 1360
    source: 11.0.0.1, destination: 13.0.0.1, id: 0xEF7D, ttl: 63,
    TOS: 0 prot: 17, source port 48823, destination port 5004

    (depth/weight/total drops/no-buffer drops/interleaves) 55/32384/7187/0/0
    Conversation 28, linktype: ip, length: 1032
    source: 12.0.0.1, destination: 13.0.0.1, id: 0x39A8, ttl: 61, prot: 1
```

Non compliant (token bucket) part from reserved flow. Weight 32,384 in WFQ. Best-effort


Configuration and monitoring

Example 1. Throughput measurements without RSVP



Configuration and monitoring

Example 1. Throughput measurements with RSVP



Configuration and monitoring

Example 2. FF style reservations

```
while true;
do packetHcli -i tap0 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/udp-size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 185;
sleep 1; done
```

ip rsvp reservation-host 13.0.0.1 11.0.0.1 udp 5004 48823 ff load 500 10
ip rsvp reservation-host 13.0.0.1 12.0.0.1 udp 5004 48823 ff load 500 10


```
ip rsvp sender 13.0.0.1 11.0.0.1 udp 5004 48823 11.0.0.1 FastEthernet0/0 1000 10  
ip rsvp sender 13.0.0.1 12.0.0.1 udp 5004 48823 11.0.0.1 FastEthernet0/0 1000 10
```

```
packetHcli -i tap1 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/udp-tap1_size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 0
```

- Two FF reservations (source port must be included)
- 1000 Kbps asked for and 500 Kbps allowed

Configuration and monitoring

Example 2. FF style reservations. Throughput measurements



Configuration and monitoring

Reservation styles by example

Ask for BW from R2

```
R2(config)# ip rsvp sender 13.0.0.1 11.0.0.1 UDP 5004 48823 11.0.0.1 FastEthernet0/0 1000 10
R2(config)# ip rsvp sender 13.0.0.1 12.0.0.1 UDP 5004 48823 12.0.0.1 FastEthernet0/1 1000 10
```

FF BW allocation from R3

```
R3(config)# ip rsvp reservation-host 13.0.0.1 12.0.0.1 UDP 5004 48823 FF LOAD 500 10
R3(config)# ip rsvp reservation-host 13.0.0.1 11.0.0.1 UDP 5004 48823 FF LOAD 500 10
```

Reserved resources at R2

```
R2# show ip rsvp reservation
```

To	From	Pro	DPort	Sport	Next Hop	I/F	Fi	Serv	BPS
13.0.0.1	11.0.0.1	UDP	5004	48823	172.16.0.2	Se1/0	FF	LOAD	500K
13.0.0.1	12.0.0.1	UDP	5004	48823	172.16.0.2	Se1/0	FF	LOAD	500K



Configuration and monitoring

Reservation styles by example

Reserve more than available

```
R3(config)# no ip rsvp reservation-host 13.0.0.1 12.0.0.1 UDP 5004 48823 FF LOAD 500 10
R3(config)# no ip rsvp reservation-host 13.0.0.1 11.0.0.1 UDP 5004 48823 FF LOAD 500 10

R3(config)# ip rsvp reservation-host 13.0.0.1 11.0.0.1 UDP 5004 48823 FF LOAD 800 10

R2# show ip rsvp reservation
To          From          Pro DPort Sport Next Hop      I/F      Fi Serv BPS
13.0.0.1    11.0.0.1          UDP 5004  48823 172.16.0.2  Se1/0    FF LOAD 800K

R3(config)# ip rsvp reservation-host 13.0.0.1 12.0.0.1 UDP 5004 48823 FF LOAD 800 10

R2# show ip rsvp reservation
To          From          Pro DPort Sport Next Hop      I/F      Fi Serv BPS
13.0.0.1    11.0.0.1          UDP 5004  48823 172.16.0.2  Se1/0    FF LOAD 800K
```

An error message returned from R4 telling that not enough BW



Configuration and monitoring

Reservation styles by example

SE style reservations

```
R3(config)# no ip rsvp reservation-host 13.0.0.1 11.0.0.1 UDP 5004 48823 FF LOAD 800 10
R3(config)# no ip rsvp reservation-host 13.0.0.1 12.0.0.1 UDP 5004 48823 FF LOAD 800 10

R3(config)# ip rsvp reservation-host 13.0.0.1 11.0.0.1 UDP 5004 48823 SE LOAD 500 10
R3(config)# ip rsvp reservation-host 13.0.0.1 12.0.0.1 UDP 5004 48823 SE LOAD 800 10
```

```
R2# show ip rsvp reservation
```

To	From	Pro	DPort	Sport	Next Hop	I/F	Fi	Serv	BPS
13.0.0.1	11.0.0.1	UDP	5004	48823	172.16.0.2	Se1/0	SE	LOAD	500K
13.0.0.1	12.0.0.1	UDP	5004	48823	172.16.0.2	Se1/0	SE	LOAD	800K

```
R2#show ip rsvp installed
```

```
RSVP: Serial11/0
```

BPS	To	From	Protoc	DPort	Sport	Weight	Conversation
800K	13.0.0.1	11.0.0.1	UDP	5004	48823	6	265

Not true
Look at queueing

Maximum shared (800 Kbps) for both sessions

Configuration and monitoring

Reservation styles by example

SE style reservations

```
R3(config)# no ip rsvp reservation-host 13.0.0.1 11.0.0.1 UDP 5004 48823 FF LOAD 800 10
R3(config)# no ip rsvp reservation-host 13.0.0.1 12.0.0.1 UDP 5004 48823 FF LOAD 800 10

R3(config)# ip rsvp reservation-host 13.0.0.1 11.0.0.1 UDP 5004 48823 SE LOAD 500 10
R3(config)# ip rsvp reservation-host 13.0.0.1 12.0.0.1 UDP 5004 48823 SE LOAD 800 10
```

```
R2# show ip rsvp reservation
```

To	From	Pro	DPort	Sport	Next Hop	I/F	Fi	Serv	BPS
13.0.0.1	11.0.0.1	UDP	5004	48823	172.16.0.2	Se1/0	SE	LOAD	500K
13.0.0.1	12.0.0.1	UDP	5004	48823	172.16.0.2	Se1/0	SE	LOAD	800K

```
R2#show ip rsvp installed
```

```
RSVP: Serial1/0
```

BPS	To	From	Protoc	DPort	Sport	Weight	Conversation
800K	13.0.0.1	11.0.0.1	UDP	5004	48823	6	265

Not true
Look at queueing

```
R2#show queueing interface Serial 1/0
Interface Serial1/0 queueing strategy: fair
```

```
(depth/weight/total drops/no-buffer drops/interleaves) 2/6/0/0/0
Conversation 265, linktype: ip, length: 1360
source: 12.0.0.1, destination: 13.0.0.1, id: 0xEF7D, ttl: 63,
TOS: 0 prot: 17, source port 48823, destination port 5004
```

Configuration and monitoring

Reservation styles by example

WF style reservations

```
R3(config)# no ip reservation-host 13.0.0.1 11.0.0.1 UDP 5004 48823 SE LOAD 500 10
R3(config)# no ip reservation-host 13.0.0.1 12.0.0.1 UDP 5004 48823 SE LOAD 800 10

R3(config)# ip rsvp reservation-host 13.0.0.1 11.0.0.1 UDP 5004 48823 WF LOAD 500 10
R3(config)# ip rsvp reservation-host 13.0.0.1 12.0.0.1 UDP 5004 48823 WF LOAD 800 10

R2# show ip rsvp reservation
To          From          Pro DPort Sport Next Hop      I/F      Fi Serv BPS
13.0.0.1    0.0.0.0          UDP 5004  0      172.16.0.2  Se1/0    WF LOAD 800K
```

Wildcards (IP 0.0.0.0 and port 0) for any source. Maximum reservation used (800 Kbps)




Configuration and monitoring

Example 3. Guaranteed-rate. Using PQ

```
while true;
do packETHcli -i tap0 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/udp-size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 185;
sleep 1; done
```

```
ip rsvp reservation-host 13.0.0.1 11.0.0.1 udp 5004 48823 ff load 400 5
ip rsvp reservation-host 13.0.0.1 12.0.0.1 udp 5004 48823 ff load 500 10
```



```
ip rsvp sender 13.0.0.1 11.0.0.1 udp 5004 48823 11.0.0.1 FastEthernet0/0 500 5
ip rsvp sender 13.0.0.1 12.0.0.1 udp 5004 48823 12.0.0.1 FastEthernet0/1 1000 10
ip rsvp pq-profile 62500 5000 ignore-peak-value (62500 Bps -> 500 Kbps , 5000 Bytes)
```

```
packETHcli -i tap1 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/udp-tap1_size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 0
```


- Traffic from `tap0` considered priority. Reserved rate (r) 400 Kbps, burst size (b) 5 KB
- Traffic from `tap1`. Reserved 500 Kbps, burst size 10 KB

Configuration and monitoring

Example 3. Guaranteed-rate. Using PQ

```
while true;
do packETHcli -i tap0 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/udp-size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 185;
sleep 1; done
```

```
ip rsvp reservation-host 13.0.0.1 11.0.0.1 udp 5004 48823 ff load 400 5
ip rsvp reservation-host 13.0.0.1 12.0.0.1 udp 5004 48823 ff load 500 10
```



```
ip rsvp sender 13.0.0.1 11.0.0.1 udp 5004 48823 11.0.0.1 FastEthernet0/0 500 5
ip rsvp sender 13.0.0.1 12.0.0.1 udp 5004 48823 12.0.0.1 FastEthernet0/1 1000 10
ip rsvp pq-profile 62500 5000 ignore-peak-value (62500 Bps -> 500 Kbps , 5000 Bytes)
```

```
packETHcli -i tap1 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/udp-tap1_size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 0
```

R2 configuration for PQ

```
ip rsvp pq-profile 62500 5000 ignore-peak-value
```

Max. rate (r') in Bps
62500 Bps = 500 Kbps

Max. burst (b') in Bytes
5000 B = 5 KB




Configuration and monitoring

Example 3. Guaranteed-rate. Using PQ

```
while true;
do packetHcli -i tap0 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/udp-size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 185;
sleep 1; done
```

```
ip rsvp reservation-host 13.0.0.1 11.0.0.1 udp 5004 48823 ff load 400 5
ip rsvp reservation-host 13.0.0.1 12.0.0.1 udp 5004 48823 ff load 500 10
```



```
ip rsvp sender 13.0.0.1 11.0.0.1 udp 5004 48823 11.0.0.1 FastEthernet0/0 500 5
ip rsvp sender 13.0.0.1 12.0.0.1 udp 5004 48823 12.0.0.1 FastEthernet0/1 1000 10
ip rsvp pq-profile 62500 5000 ignore-peak-value (62500 Bps -> 500 Kbps , 5000 Bytes)
```

```
packetHcli -i tap1 -d 6000 -m 2 -f /home/cesar/GNS3/Captures/udp-tap1_size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 0
```

Every reservation such that:

$$r \leq r' \text{ and } b \leq b'$$


will be considered into PQ

Configuration and monitoring

Example 3. Guaranteed-rate. Using PQ

```
while true;
do packETHcli -i tap0 -d 6000 -m 2 -f /home/cesar/GN53/Captures/udp-size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 185;
sleep 1; done
```

```
ip rsvp reservation-host 13.0.0.1 11.0.0.1 udp 5004 48823 ff load 400 5
ip rsvp reservation-host 13.0.0.1 12.0.0.1 udp 5004 48823 ff load 500 10
```



```
ip rsvp sender 13.0.0.1 11.0.0.1 udp 5004 48823 11.0.0.1 FastEthernet0/0 500 5
ip rsvp sender 13.0.0.1 12.0.0.1 udp 5004 48823 12.0.0.1 FastEthernet0/1 1000 10
ip rsvp pq-profile 62500 5000 ignore-peak-value (62500 Bps -> 500 Kbps , 5000 Bytes )
```

```
packETHcli -i tap1 -d 6000 -m 2 -f /home/cesar/GN53/Captures/udp-tap1_size_1356-port_dest_5004_ip_13.0.0.1.pcap -n 0
```

Monitoring reservations

```
R2#show ip rsvp installed
```

```
RSVP: FastEthernet0/0 has no installed reservations
```

```
RSVP: Serial1/0
```

BPS	To	From	Protoc	DPort	Sport	Weight	Conversation
400K	13.0.0.1	11.0.0.1	UDP	5004	48823	0	264
500K	13.0.0.1	12.0.0.1	UDP	5004	48823	6	265

```
RSVP: FastEthernet0/1 has no installed reservations
```


Contents

- 1 QoS Overview
- 2 Classification
- 3 Congestion Management
- 4 Congestion Avoidance
- 5 Policing and shaping
- 6 Resource Reservation Protocol
- 7 Bibliography**



Bibliography

- *Network Warrior 2nd Ed.* Gary A. Donahue. O'Reilly, 2011
- Cisco IOS Quality of Service Solutions Configuration Guide, Release 12.2SR
- Cisco. Quality of Service Networking
- Cisco IOS Quality of Service Solutions Command Reference
- *Administering CISCO IP QoS in IP Networks* . Syngress, 2001
- *Internet QoS: Architectures and Mechanisms for Quality of Service*. Zheng Wang. Morgan Kaufmann Publishers, 2001
- Cisco. QoS: RSVP Configuration Guide

